


Raw Attention from the Answer Position in Unedited LLaVA-v1.5-7B



Top: raw head-averaged attention from the answer position to text tokens across all 32 layers. Bottom: selected-layer attention to visual patches, normalized per layer for display.